

Beyond the Chatbot

Grounding AI in the Law

John A. Kelly

Office of the County Counsel, County of Sonoma

July 20, 2026

THE HOUR AHEAD

STATUTES

STRAIGHTFORWARD

LOCAL RULES

COLLECTED

CASE LAW

PAYWALLED

**YOUR OWN
DOCUMENTS**


OVERLOOKED

blog.law-kelly.com | MCLE

A Process Problem, Not a Technology Problem


The electric motor took a generation to pay off – because factories kept the old layout

1890s



New tool, same layout – productivity barely moved.

1920s



New process – transformed results.

THE SAME CHOICE, TODAY



Open chatbot.
Paste question.
File what comes back.

NEW TOOL, OLD PROCESS



Ground.
Verify.
Document.

NEW PROCESS

AI hallucination is a consequence of unchanged process.
Fix the process and you fix the hallucination.

The Model Needs a Harness

Even AI's leading skeptic says the value is the scaffolding around the chatbot — not the chatbot

WHO — AND WHY TRUST HIM

Gary Marcus

- ✓ Cognitive scientist; professor emeritus of psychology and neuroscience, NYU
- ✓ Founded an AI company (Geometric Intelligence) — acquired by Uber
- ✓ Warned LLMs would hallucinate years before ChatGPT — and was right
- ✓ A builder and an early-right skeptic — not selling the hype

"They call it the harness, but that harness is really symbolic AI."

"The neural networks are pretty bad at sticking to facts. So that's the hallucination stuff ... and the symbolic stuff, they don't make stuff up that way."

— Gary Marcus, in conversation with Brian Greene, World Science Festival (May 2026)

STRUCTURAL BLIND SPOTS — IT PREDICTS TEXT, IT DOES NOT REASON

Not a chess engine — a 1979 Atari 2600 chess program beat ChatGPT (2025): illegal moves, lost the board

Not a calculator — without a tool, it only estimates the product of two long numbers

Not a database — it returns what sounds right, not what is verified true

The fix is not a smarter chatbot. It is the harness around it — grounding, verification, documentation.


Where Courts and Practitioners Stand

6 L.A. judges in the Learned Hand pilot


10 states using the program

17-33% hallucination rate in commercial legal AI (Stanford 2025)

ON THE BENCH



AT COUNSEL TABLE — THE SANCTIONS TIMELINE



You are not deciding whether to use AI. You are already responding to AI work product.

Five Duties, New Failure Modes

Not new obligations — new ways to violate them

COMPETENCE

Model Rule 1.1, Comment 8

Knowing the failure rates **IS** competence.

CANDOR

Bus. & Prof. Code § 6068(d) | Rule 3.3

A fabricated citation in a filing **cannot be sat on**.

DISCLOSURE

Courts increasingly require disclosure of AI assistance. Direction clear, **not yet uniform**.

VERIFICATION

Output looks correct by design.
The check must be mechanical, **not a smell test**.

SUPERVISION

Rules 5.1 / 5.3

Staff are already using these tools. Written protocol, **not a hallway conversation**.

THE STANDARD OF CARE IS MOVING

Electronic research — standard of care, ~2010

AI — the same jump now, 2026

Once a tool becomes standard, competent use of it is the baseline — and for AI, competent use means verification. Opponents already stress-test filings with AI.

The fabrication was the model's. The filing is yours.

Four Types of Authority

AI hallucinates in all four — the discipline is the same

STATUTES

**29 codes,
bulk-downloadable**

easiest to ground



LOCAL RULES

**58 counties,
96 PDFs**

now collected —
available to you



CASE LAW

**paywalled,
no bulk access**

collectible with effort,
case by case



RAW MATERIALS

**emails, contracts,
transcripts**

hardest to verify —
most hallucination



EASIEST TO GROUND

HARDEST TO VERIFY

Same discipline for all four: ground the AI in sources, verify mechanically.

Statutory Hallucination and the Grounding Principle

UNGROUND

Training data — millions
of statutory citations

pattern
matching

Civ. Code ✕
§ 1542.7(b)(3)

Plausible format — but § 1542 has no
subdivisions. Invented; does not exist.

17-33%

hallucination rate,
ungrounded commercial tools

GROUND

Actual code text —
Legislature bulk data

grounding

Civ. Code § 1542 ✓

The real general-release statute.
Verified against source.

~0%

grounded in the actual text

Magesh et al., Stanford 2025. Butler & Butler
2026: retrieval quality is the primary driver.

Fix the inputs. The outputs follow.

The California Codes Pipeline

29

California codes

9.6 MB

plain text

nightly

sync from official bulk data

**Legislature bulk
data archive**

downloads.leginfo.
legislature.ca.gov
— ~764 MB ZIP

**Parse and
extract all 29
codes**

9.6 MB clean
plain text

**GitHub repository —
public, nightly sync**

[github.com/johnakelly-yahoo-com/
california-codes](https://github.com/johnakelly-yahoo-com/california-codes)

**AI reads the
actual statute
text**

SECTION 1. No person shall...
unless authorized by law.



the model quotes the provided text verbatim

If the authoritative text is public, download it and ground the model. You get this link today.

downloads.leginfo.legislature.ca.gov

Few-Shot Citation Prompting

Teach the model your style manual by example

WITHOUT EXAMPLES

People v. Sanchez, 63 Cal.4th 665 (2016)



Bluebook habit — year at the end, no parallel cite. Wrong manual for California briefs.

Works for any jurisdiction, any style manual.

EXAMPLES PROVIDED TO THE MODEL

Sargon Enterprises, Inc. v. USC (2012) 55 Cal.4th 747 [149 Cal.Rptr.3d 614, 288 P.3d 1237]

Lazar v. Superior Court (1996) 12 Cal.4th 631 [49 Cal.Rptr.2d 377, 909 P.2d 981]

Aguilar v. Atlantic Richfield Co. (2001) 25 Cal.4th 826 [107 Cal.Rptr.2d 841, 24 P.3d 493]

California Style Manual: year in parentheses after the names; official and parallel cites bracketed.

The parallel cite is a built-in cross-check — nothing slips through.

Provide the pattern; the model follows it.

Three examples beat three paragraphs of instructions.

WITH EXAMPLES

People v. Sanchez (2016) 63 Cal.4th 665 [204 Cal.Rptr.3d 102, 374 P.3d 320]



Show, don't tell. 3-5 examples set the format.

The 58-County Problem

58 Counties

96 PDF Files

148 MB of text

Each county writes its own local rules. They differ on law and motion, filing, discovery, case management, ADR, ex parte, e-filing, tentative rulings, trial setting, and fees. No lawyer, committee, or agency had ever synthesized all 58.

Presented as a case study in AI-assisted legal research methodology.

THE RULES DIVERGE



EXAMPLE 1: FILING DEADLINES —
County A: 4pm. County B: 5pm.
County C: noon.

EXAMPLE 2: ADR REQUIREMENTS —
mandatory mediation vs. voluntary
settlement conference vs. no rule at all.

EXAMPLE 3: EX PARTE NOTICE —
24-hour notice vs. 10am the day
prior vs. the judge's discretion.

SYNTHESIS EFFORT AND TIME

Traditional method — human review

thousands of hours (estimated)


AI-assisted — machine reading and analysis

hours to days (demonstrated)

rapid, comprehensive analysis.

The information exists. Nobody had time to read it all.

The Survey Pipeline and the Deliverable



Sonoma vs. each county,
10 categories — 43 minutes total

175,000 tokens → 5,000

57
memoranda

78
distinct practices catalogued

10
recommendations

ADOPTION TIERS

WIDESPREAD — 30+ counties
Examples: e-filing, tentative rulings, vertical assignment

MODERATE — 14-22 counties
Examples: exhibit bookmarking, disposition tracks

INNOVATIVE — 6-9 counties
Examples: online reservation with forfeiture, informal discovery conference

NOVEL — 5 or fewer
Examples: no case off calendar, fee schedule

The pipeline handled the scope. The attorney supplied the judgment.


What the Model Got Wrong — and What Judgment Supplied

MODEL OUTPUT		AFTER ATTORNEY REVIEW	WHAT AI CONTRIBUTED	WHERE JUDGMENT WAS INDISPENSABLE
Sonoma lacks tentative rulings	→	Local rule 5.5 operates as auto-adoption	● Scope: all 58 jurisdictions	● Accuracy: correcting factual errors
6 counties = widely adopted	→	6 counties = innovative (precise tier)	● Structure: 10 categories, consistent format	● Preemption: Cal. Rules of Ct., rule 3.20 analysis
No preemption analysis	→	Full preemption assessment, Cal. Rules of Ct., rule 3.20	● Coverage: 78 distinct practices	● Interpretation: existing-practice context
			● Speed: 43 minutes for 57 comparisons	● Terminology: field-specific precision

Comprehensive coverage is not accuracy. Let the AI handle breadth; supply the judgment; document the corrections.

The Smallest Case: A \$301 Red Light Ticket

The same discipline, end to end — worked as a case study in your written materials



Not guilty on the written declarations. Bail refunded in full.

CASE LAW IS DIFFERENT

STATUTES ✓

bulk download available — ground the model directly

CASE LAW ✗

paywalled, no bulk access — you cannot ground what you cannot download

THREE FAILURE MODES

Fabricated case — does not exist

✓ easy to catch: a 10-second lookup

The model invents a case name and citation (e.g., “*Smith v. Jones*, 123 F.3d 456 (2023)”). A simple database search confirms it is fake.

Real case — fabricated quotation

⚠ moderate: full-text search finds it

The model correctly cites a real case (e.g., “*Brown v. Board of Education*”), but includes a plausible-sounding quotation that does not appear in the actual opinion. Full text analysis reveals the mismatch.

Real case — mischaracterized holding

⚠ dangerous: passes the quick checks

the case is real, the quote is real, but it does not stand for the proposition cited. The model uses a dissenting opinion or dictum as the binding holding, misleading the user on the law.

The third failure mode is the one that gets through.

The Verification Ladder

Four steps, increasing rigor — each catches what the last one missed

1

Does the case exist?

Westlaw or Lexis pull — 10 seconds

every sanctioned
attorney failed
HERE

2

Does the quotation exist?

full-text search of the opinion — 30 seconds

3

Good law for THIS proposition?

headnotes and Key Numbers — the key step

4

Still good law?

KeyCite or Shepard's on the specific point

Opinion — 30 pages

[1]

[2]

[3]

[4]

HEADNOTES

[1] Negligence:
Duty of Care

[2] Evidence:
Expert Testimony

[3] Damages:
Economic Loss

[4] Procedure:
Summary Judgment

AI says the case is about economic loss →
headnote [3] confirms → jump straight to
the passage. Seconds, not thirty pages.

Each step catches what the previous one missed.

Raw Materials — Where the Real Hallucination Happens

Emails

Correspondence

Contracts

Financial Records

AI can fabricate deposition quotes, invent contract clauses, misattribute statements.

Deposition
Transcripts

Expert Reports


Discovery
Responses

Mediation Briefs

WHAT IS A LINTER?


Borrowed from software: a tool that automatically checks work against a fixed rule and flags every violation. Here the rule is simple — every quote and citation must appear in the source.

LEVEL 1 — STRING MATCH



purely mechanical — catches fabrication

LEVEL 2 — BACKED BY A SECOND AI



a second AI reads the cited passage and judges whether the characterization matches — catches what string-matching cannot

Most real-world hallucination happens here — and nobody is talking about it.

The Treatise and the Linter

Provide the treatise — Rutter Guide, Witkin. Instruct: cite ONLY authorities appearing in this source. Anything outside the corpus is, by definition, a hallucination.

AI OUTPUT (illustrative)

"...the duty of candor..." — p. 212

"...reasonable inquiry standard..." — p. 88

"...competence includes technology..." — p. 305

"...sanctions are mandatory here..." — p. 401



LINTER

string comparison against the source text

candor

duty of candor [p. 212]

inquiry

reasonable inquiry standard [p. 88]

technology

competence includes technology [p. 305]

sanctions

sanctions are mandatory here [p. 401]

exact string-match verification



RESULTS



found in the treatise



found in the treatise



found in the treatise



flagged: not in the treatise

Mechanical. Automatic. Runs before any substantive review.

Not subjective. A string comparison. Yes or no.

Structured Prompts for Westlaw and Lexis

Integration, not replacement

THE VERIFICATION PROMPT

For each citation:

1. confirm existence
2. locate the quotation
3. check headnotes for the proposition
4. flag for citator review

1. Confirm existence

2. Locate the quotation

3. Check headnotes

the key
step

4. Flag for citator review



Verified citation —
ready to rely on




Flagged — attorney
review required

The AI does not replace Westlaw or Lexis. It structures your workflow for using them. Tested prompts are in your handout materials.

The model handles the repetitive checking. You make the legal judgments.

The Harness — Full Architecture

The redesigned factory, in three layers



Six Things This Week

1

Stress-test your next brief with AI before filing. Your opponent will.

Pre-filing analysis for weaknesses and counterarguments.

2

Verify every AI citation: existence, quotation, headnotes, citator.

Confirm accuracy and validity against primary sources.

3

Provide source material. Never rely on training data alone.

Upload the relevant cases, statutes, or documents.

4

Professional tier. Training off. Treat sessions as work product.

Keep client information out of the model's training set.

30 seconds
to turn off training
on your inputs

5

Start a work log: date, task, result, details.

Maintain a clear audit trail of AI use.

6

Write the staff verification protocol. An SOP, not a habit.

Codify human review of AI-generated content.

Concrete. Actionable. Start here.

What You Take Home

Three electronic resources — links live now

All 29 California Codes

machine-readable, nightly sync from the Legislature's bulk data

github.com/johnakelly-yahoo-com/california-codes

Usage: Point the model at the Probate Code file and instruct: answer only from this text.

All 58 County Local Rules

structured for AI retrieval — full rules plus civil-only extracts

github.com/johnakelly-yahoo-com/california-local-rules

Usage: Load your county's file before drafting anything the local rules govern.

Tested Prompts

case law verification, citation checking, document-grounded research — designed for Westlaw and Lexis

in your handout materials

Usage: Run the four-step verification prompt on every AI-drafted brief.

Links live now. Implement this week.

“Beyond the Chatbot”

Grounding AI in the Law



Ground it in the sources.



Verify it mechanically.



Document what you did.

The method of finding the law shapes what law you find.

See generally Robert C. Berring, *Legal Research and Legal Concepts: Where Form Molds Substance* (1987) 75 Cal. L.Rev. 15